

## Data Collection and Preprocessing Phase

Date	15 March 2024
Team ID	LTVIP2024TMID24981
Project Title	Deep learning techniques for breast cancer prediction
Maximum Marks	6 Marks

### Data Exploration and Preprocessing Report

A **Data Exploration and Preprocessing Report** is a crucial step in developing a robust model for **breast cancer prediction**. The goal is to thoroughly explore the dataset, understand its structure, and apply the necessary preprocessing steps before feeding the data into a machine learning model such as a CNN. Below is an outline of what such a report would include, tailored to a breast cancer prediction project using images (e.g., mammograms or histopathology slides).

Section	Description
Data Overview	<p>Breast cancer prediction using deep learning is an active research area aimed at improving early detection, diagnosis, and prognosis of breast cancer. Deep learning methods, particularly convolutional neural networks (CNNs).</p> <ul style="list-style-type: none"> <li>• <b>Dataset Source:</b> Identify the source of the dataset (e.g., publicly available datasets like <b>Break His</b>, <b>DDSM</b>, or a hospital dataset).</li> <li>• <b>Data Type:</b> Describe the type of data (e.g., mammograms, histopathology images, or both). <ul style="list-style-type: none"> <li>◦ <b>Mammograms:</b> Typically, grayscale images used to detect masses or calcifications.</li> <li>◦ <b>Histopathology Slides:</b> Usually RGB-stained images used to analyse cellular structures.</li> </ul> </li> <li>• <b>Dataset Size:</b> Provide the number of samples, including breakdowns by class (e.g., benign vs. malignant).</li> <li>• <b>Labels:</b> Outline the labels and their distribution. Common labels might include: <ul style="list-style-type: none"> <li>◦ <b>Benign</b></li> <li>◦ <b>Malignant</b></li> <li>◦ <b>Normal</b></li> </ul> </li> </ul> <p><b>1.2. Class Distribution</b></p>

	<ul style="list-style-type: none"> <li>• <b>Class Imbalance:</b> Explore the distribution of classes (benign, malignant, normal). If there's a significant imbalance (e.g., far more benign than malignant cases), consider strategies like oversampling, under sampling, or synthetic data generation (e.g., SMOTE).</li> </ul> <p><b>1.3. Image Dimensions</b></p> <ul style="list-style-type: none"> <li>• <b>Image Sizes:</b> Check the dimensions of the images. Often, images from different sources may have varying resolutions. This impacts how we process and resize them.</li> <li>• <b>Colour Channels:</b> Investigate whether the images are grayscale or RGB. Mammograms are usually grayscale, while histopathology slides are RGB.</li> </ul> <p><b>1.4. Missing Data</b></p> <ul style="list-style-type: none"> <li>• <b>Check for Missing Images or Labels:</b> Ensure that all images have corresponding labels and there are no corrupted files in the dataset.</li> </ul> <p><b>1.5. Visual Inspection</b></p> <ul style="list-style-type: none"> <li>• <b>Sample Images:</b> Display a few random images from each class (benign, malignant) to gain an intuitive understanding of the dataset. This helps in understanding image quality and variability.</li> </ul>
Resizing	<ul style="list-style-type: none"> <li>• <b>Objective:</b> Standardize the input image size for the model.</li> <li>• <b>Why Resize:</b> Different imaging devices can produce images of varying dimensions. Neural networks typically require all inputs to have a fixed size. For instance, you may resize all images to 224x224 for compatibility with models like Resnet or Efficient Net.</li> <li>• <b>Consideration:</b> While resizing, preserving the aspect ratio is crucial to avoid distortion of important features like cell structures or tumour boundaries.</li> </ul>
Normalization	<ul style="list-style-type: none"> <li>• <b>Objective:</b> Ensure that the pixel intensity values are on a similar scale for better training stability.</li> <li>• <b>Why Normalize:</b> Medical images (like histopathology or mammograms) can have varying intensity ranges due to differences in scanners, staining techniques, or patient conditions. Normalizing the pixel values (scaling to [0, 1] or [-1, 1]) ensures that the learning algorithm treats all features uniformly.</li> <li>• <b>Techniques:</b> Common normalization approaches include subtracting the mean and dividing by the</li> </ul>

	standard deviation for pixel intensities.
Data Augmentation	<p><b>Objective:</b> Artificially increase the dataset size to make the model more robust.</p> <p><b>Why Augment:</b> Medical imaging datasets are often small and imbalanced. Data augmentation can help mitigate this by applying transformations such as:</p> <ul style="list-style-type: none"> <li>• <b>Rotations:</b> Slightly rotating images can help the model recognize tumours regardless of the orientation.</li> <li>• <b>Flips:</b> Horizontal and vertical flips can help the model learn symmetries.</li> <li>• <b>Zoom:</b> Zooming in or out simulates variability in zoom levels between different images.</li> <li>• <b>Contrast Adjustments:</b> Varying contrast can help simulate different imaging conditions.</li> </ul>
Denoising	<ul style="list-style-type: none"> <li>• <b>Objective:</b> Reduce noise to improve image quality and model performance.</li> <li>• <b>Why Denoise:</b> Medical images, especially histopathological images, can have noise due to artifacts from the imaging process, patient movement, or machine errors.</li> <li>• <b>Techniques:</b> <ul style="list-style-type: none"> <li>○ <b>Gaussian Blur:</b> Reduces high-frequency noise, which may not be relevant for feature extraction.</li> <li>○ <b>Median Filtering:</b> Effective for removing 'salt-and-pepper' noise that might appear due to image acquisition errors.</li> <li>○ <b>Deep Learning Methods:</b> Denoising autoencoders or other deep learning-based techniques can be trained to remove noise while preserving important features like tumour boundaries.</li> </ul> </li> </ul>
Edge Detection	<p><b>Objective:</b> Highlight the boundaries of tumours or other regions of interest.</p> <p><b>Why Edge Detection:</b> Tumour boundaries and irregular shapes are important in cancer prediction models. Edge detection helps isolate these features for analysis.</p> <p><b>Techniques:</b></p> <ul style="list-style-type: none"> <li>• <b>Sobel Filter:</b> Simple technique for detecting edges based on intensity gradients.</li> </ul>

	<ul style="list-style-type: none"> <li>• <b>Canny Edge Detection:</b> More sophisticated, useful for detecting sharp changes in intensity which often correspond to tumour borders.</li> <li>• <b>Application:</b> Edge detection can help segment regions of interest (e.g., tumours) from surrounding tissues, improving the focus on important features.</li> </ul>
Color Space Conversion	<p><b>Objective:</b> Convert images from one colour space (like RGB) to another (such as grayscale or HSV) to highlight specific features useful for prediction.</p> <p><b>Why Convert Colour Spaces:</b></p> <ul style="list-style-type: none"> <li>• <b>Histopathology Images:</b> In breast cancer prediction, histopathological images are often stained using techniques like H&amp;E (hematoxylin and eosin). Colour information in these images can be crucial for identifying cancerous tissues.</li> <li>• <b>Mammograms:</b> Mammograms are usually grayscale images, and colour space conversion may not be necessary. However, processing them in different intensity ranges can help highlight contrasts in the tissue.</li> </ul>
Image Cropping	<ul style="list-style-type: none"> <li>• <b>Objective:</b> Crop out unnecessary parts of the image to focus on the region of interest (ROI), such as a tumour.</li> <li>• <b>Why Crop:</b> Medical images often contain a lot of background information that may not be relevant for the model. Cropping helps: <ul style="list-style-type: none"> <li>○ <b>Remove Background Noise:</b> Focus on key areas where tumours or abnormalities are likely to exist.</li> <li>○ <b>Reduce Image Size:</b> Cropping can reduce image dimensions, which speeds up model training and reduces computational load.</li> <li>○ <b>Improves Model Performance:</b> The CNN will focus on the important parts of the image (e.g., suspicious regions in a mammogram or biopsy), leading to better learning and prediction accuracy.</li> </ul> </li> </ul>
Batch Normalization	<p><b>Objective:</b> Normalize the input features across a mini-batch, improving training speed and model stability.</p> <p><b>Why Use Batch Normalization?:</b> In deep learning, the internal covariate shift (changing data distributions during training) can</p>

	<p>slow down learning. Batch normalization helps by:</p> <ul style="list-style-type: none"> <li>• <b>Stabilizing Learning:</b> Normalizes activations layer by layer, ensuring that the distribution of inputs to each layer stays consistent throughout training.</li> <li>• <b>Faster Convergence:</b> By maintaining stable gradients, batch normalization enables faster and more efficient training of deep CNNs.</li> <li>• <b>Regularization:</b> It acts as a form of regularization, reducing overfitting by introducing noise to the network's activations.</li> </ul>
<b>Data Preprocessing Code Screenshots</b>	
Loading Data	<p>This step involves loading the image data and resizing it to a fixed size (e.g., 224x224).</p> <pre>#function to load differnet labels found in dataset path = "Dataset" labels = [] X = [] Y = [] for root, dirs, directory in os.walk(path):     for j in range(len(directory)):         name = os.path.basename(root)         if name not in labels:             labels.append(name.strip()) print(labels)</pre> <p>['benign', 'malignant', 'normal']</p>

<p>Normalization</p>	<p>Scaling the pixel values between 0 and 1.</p> <pre> ]: #preprocess images like shuffling and normalization X = X.astype('float32') X = X/255 indices = np.arange(X.shape[0]) np.random.shuffle(indices)#shuffle all images X = X[indices] Y = Y[indices] Y = to_categorical(Y) #split dataset into train and test X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2) print("Dataset Image Processing &amp; Normalization Completed") print("80% images used to train CNN algorithm : "+str(X_train.shape[0])) print("20% image used to train CNN algorithm : "+str(X_test.shape[0])) </pre> <p>Dataset Image Processing &amp; Normalization Completed 80% images used to train CNN algorithm : 1248 20% image used to train CNN algorithm : 312</p>
<p>Data Augmentation</p>	<p>Applying random transformations to increase dataset diversity.</p> <pre> ): #finding &amp; plotting graph of non-addicted and addicted instances #visualizing class labels count found in dataset label, count = np.unique(Y, return_counts = True) print("Benign : "+str(count[0])) print("Malignant : "+str(count[1])) print("Normal : "+str(count[2])) height = count bars = labels y_pos = np.arange(len(bars)) plt.figure(figsize = (4, 3)) plt.bar(y_pos, height) plt.xticks(y_pos, bars) plt.xlabel("Dataset Class Label Graph") plt.ylabel("Count") plt.xticks() plt.show() </pre> <p>Benign : 874 Malignant : 420 Normal : 266</p>

<p>Batch Normalization</p>	<pre>]: #preprocess images like shuffling and normalization X = X.astype('float32') X = X/255 indices = np.arange(X.shape[0]) np.random.shuffle(indices)#shuffle all images X = X[indices] Y = Y[indices] Y = to_categorical(Y) #split dataset into train and test X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.2) print("Dataset Image Processing &amp; Normalization Completed") print("80% images used to train CNN algorithm : "+str(X_train.shape[0])) print("20% image used to train CNN algorithm : "+str(X_test.shape[0]))</pre> <p>Dataset Image Processing &amp; Normalization Completed 80% images used to train CNN algorithm : 1248 20% image used to train CNN algorithm : 312</p>
<p>CNN Algorithm</p>	<pre>In [93]: #creating CNN object cnn_model = Sequential() #adding CNN2d layer with 32 neurons of size 3 X 3 to filter images 32 times cnn_model.add(Convolution2D(32, (3, 3), input_shape = (X_train.shape[1], X_train.shape[2], X_train.shape[3]), activation = 'relu')) #max pool layer to collect filtered relevant features from previous CNN layer cnn_model.add(MaxPooling2D(pool_size = (2, 2))) #adding another layer with relu activation function #ReLU helps the first hidden layer receive errors from the last layers to adjust all weights between layers cnn_model.add(Convolution2D(32, (3, 3), activation = 'relu')) cnn_model.add(MaxPooling2D(pool_size = (2, 2))) cnn_model.add(Flatten()) #adding LSTM as RNN layer cnn_model.add(RepeatVector(2)) cnn_model.add(LSTM(32, activation = 'relu'))#=====adding RNN LSTM #defining output layer with extra softmax layer which will divide each class prediction into probabilities and the #class with highest probability will be best prediction and help in enhancing accuracy cnn_model.add(Dense(units = 256, activation = 'relu')) cnn_model.add(Dense(units = y_train.shape[1], activation = 'softmax')) #compile the model cnn_model.compile(optimizer = 'adam', loss = 'categorical_crossentropy', metrics = ['accuracy']) #train and load the model if os.path.exists("model/cnn_weights.hdf5") == False:     model_checkpoint = ModelCheckpoint(filepath="model/cnn_weights.hdf5", verbose = 1, save_best_only = True)     hist = cnn_model.fit(X_train, y_train, batch_size = 32, epochs = 15, validation_data=(X_test, y_test), callbacks=[model_checkpoint], verbose=1)     f = open("model/cnn_history.pkl", 'wb')     pickle.dump(hist.history, f)     f.close() else:     cnn_model.load_weights("model/cnn_weights.hdf5") #perform prediction on test data using cnn model predict = cnn_model.predict(X_test) predict = np.argmax(predict, axis=1) y_test1 = np.argmax(y_test, axis=1) #call this function to true test labels and predicted labels to calculate accuracy and other metrics calculateMetrics("CNN with Softmax", y_test1, predict)</pre> <p>CNN with Softmax Accuracy : 99.35897435897436 CNN with Softmax Precision : 99.43602693602695 CNN with Softmax Recall : 99.43602693602695 CNN with Softmax FSCORE : 99.43602693602695</p>